a Settlement chain

Only the calibration function is editable

Issued forecast
P_P · 96 × 15-min intervals

→

Post-processing
transform(view) → $\hat{P}$
search can rewrite this function

→

Grid-code score
ACC from max(P_M, 0.2 Cap)

↑ metered output P_M

→

Settlement
penalised energy

b Automated search loop

60 iterations per run

Forecast view
P_P, Cap, time fields

→

Proposer
writes transform()

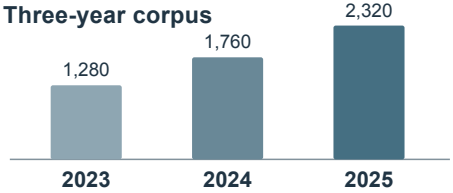
→

Execute & score
candidate on validation farms

Keep if ΔACC > τ
otherwise retain current code

One scalar score per iteration; metered output is never shown to the proposer.

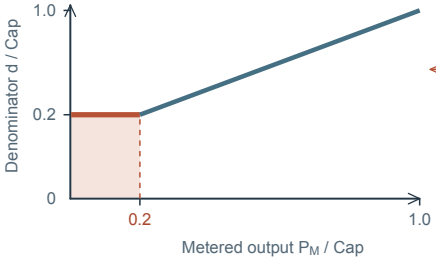
c Three-year corpus



106 runs · **5,360** iterations · **1,202** accepted edits

d Where the denominator changes

d / Cap = max(P_M / Cap, 0.2)



Below floor
fixed d
59% of intervals

Relative scaling ends; score follows distance from a fixed target.

Above floor: d = P_M

e Four tests of the apparent gain

Change one evaluation dimension at a time

1 Subset
Hold out farms or days
same metric · different evaluation subset

2 External site
Apply unchanged code
farms never used in the search

3 Independent loss
Re-score identical forecasts
capacity-normalised error · no floor

4 Regime
Split at 0.2 Cap
floor-binding vs proportional region